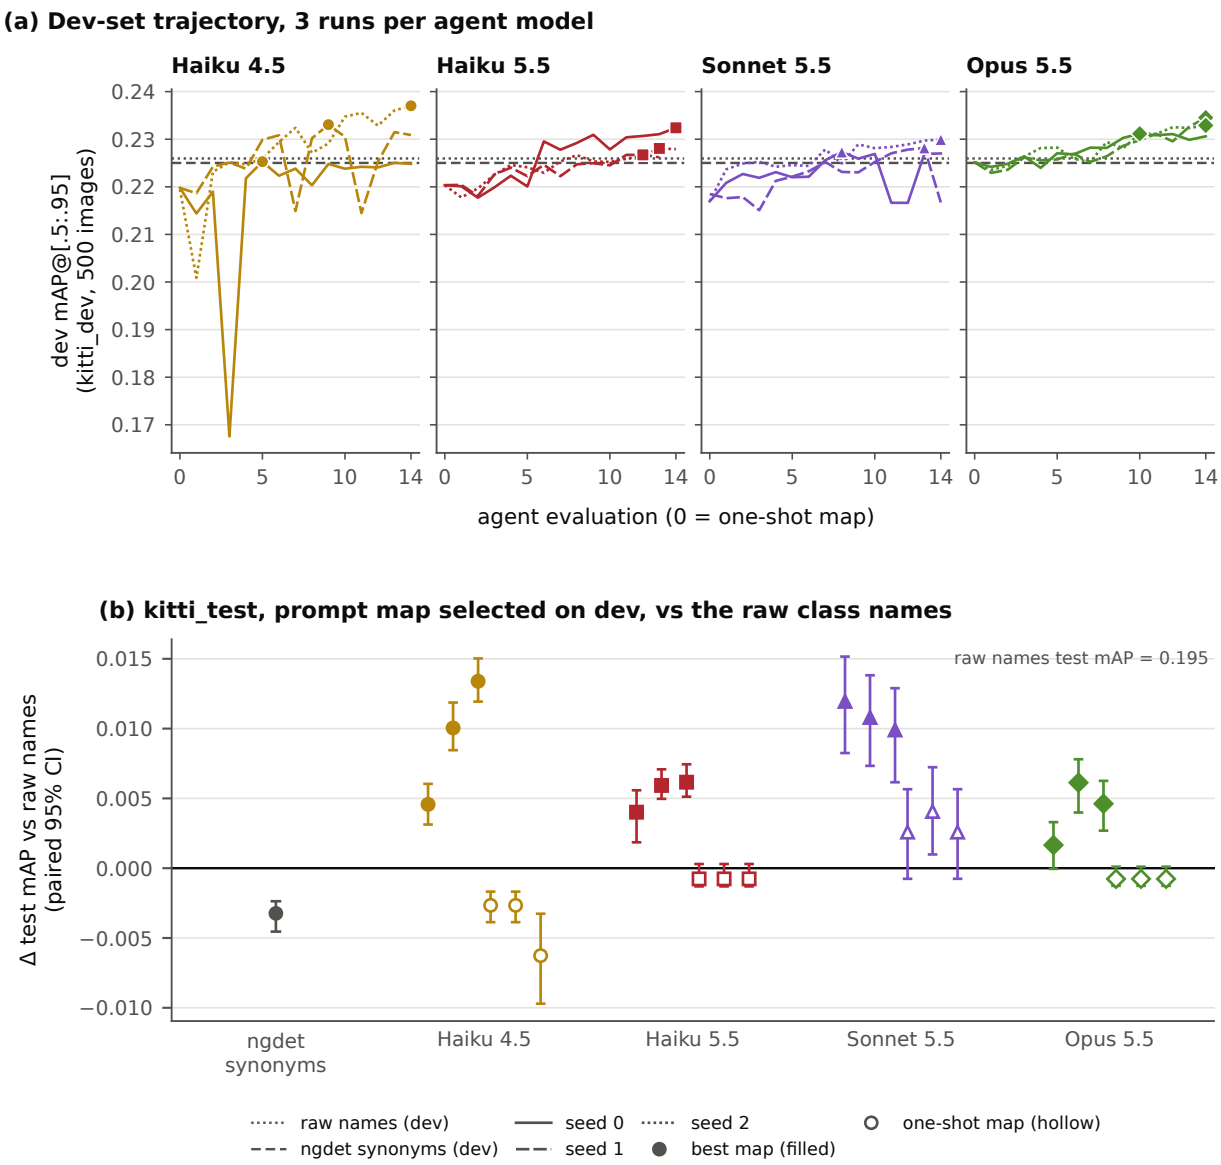

(a) Larger marker = the evaluation the agent kept as its best map. (b) Within each model: best maps of seeds 0, 1, 2, then one-shot maps of seeds 0, 1, 2. Bars are paired bootstrap CIs (1000 image resamples of kitti_test) and reflect image sampling only, not run-to-run variance of the agent. 3 runs per model. Haiku 4.5 ran without the effort setting; the other three ran at effort medium. KITTI mAP is COCO-style mAP@[.5:.95] over the 5 mapped classes (car, truck, pedestrian, cyclist, tram); COCO mAP is over 80 classes.